

CS221 Exam Solutions

Spring 2020

Read all of the following information before starting the exam:

- This test has 4 problems on 29 pages for a total of 140 possible points. **However, we will only be grading this exam out of 120 points.** Any points you get over a score of 120 will be used as extra credit.
- Note that different questions are worth different amounts of points. Budget your time accordingly!
- Keep your answers precise and concise. We may award partial credit so show all your work clearly and in order.
- Don't spend too much time on one problem. Read through all the problems carefully and do the easier ones first.
- This exam is open-book; you may use any resources, including the course website.
- Good luck!

Problem	Part	Max Score
1	a	14
	b	10
	c	6
	Total	30
2	a	4
	b	12
	c	4
	Total	20
3	a	6
	b	10
	c	12
	d	12
	Total	40
4	a	6
	b	8
	c	10
	d	4
	e	12
	f	10
	Total	50

1. Classification (30 points)

Suppose we have a dataset of N points (x_i, y_i) , where $y_i \in \{-1, +1\}$ for all i . Consider the binary classification problem on this dataset with the pointwise loss function $Loss(\mathbf{x}_i, y_i, \mathbf{w})$ defined as follows:

$$Loss(\mathbf{x}_i, y_i, \mathbf{w}) = \frac{1}{2N} \sum_{j=1}^d (w^{(j)})^2 + \lambda \max \left\{ 0, 1 - y_i \left(\sum_{j=1}^d w^{(j)} \phi(\mathbf{x}_i)^{(j)} \right) \right\} \quad \forall i \in \{1, 2, \dots, N\} \quad (1)$$

where

- $\phi(\mathbf{x}_i)$ is the feature vector corresponding to the i th data point
- $\mathbf{w} \in \mathbb{R}^d$ is the learned weight vector, with $w^{(j)}$ being its j th dimension.
- λ is a regularization parameter

The overall training loss $TrainLoss(\mathbf{w})$ is defined as follows:

$$TrainLoss(\mathbf{w}) = \sum_{i=1}^N Loss(\mathbf{x}_i, y_i, \mathbf{w})$$

a. (14 points)

Stochastic Gradient Descent

Consider the following implementation of the stochastic gradient algorithm:

Algorithm 1: Stochastic Gradient Descent

```
1: Randomly shuffle the training data
2: initialize  $\mathbf{w} = \mathbf{0}, i = 1$ 
3: for  $t = 1, 2, \dots, T$  do
4:    $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}(\mathbf{x}_i, y_i, \mathbf{w})$ 
5:    $i \leftarrow (i \bmod N) + 1$ 
6: end for
```

where η is the learning rate, t is the iteration number (different from lecture slides, where t represented the epoch number), and T is the maximum number of iterations. The rest are as defined as in Equation (1).

(i) [3 points] Compute the gradient of the pointwise loss function above with respect to \mathbf{w} , i.e., $\nabla_{\mathbf{w}} \text{Loss}(\mathbf{x}_i, y_i, \mathbf{w})$. Express your answer **in vector notation**.

Solution

$$\nabla_{\mathbf{w}} \text{Loss}(\mathbf{x}_i, y_i, \mathbf{w}) = \frac{1}{N} \mathbf{w} + \lambda \mathbf{1}[y_i(\mathbf{w} \cdot \phi(\mathbf{x}_i)) < 1](-y_i \phi(\mathbf{x}_i))$$

OR

$$\nabla_{\mathbf{w}} \text{Loss}(\mathbf{x}_i, y_i, \mathbf{w}) = \frac{1}{N} \mathbf{w} + \lambda \mathbf{1}[y_i(\mathbf{w} \cdot \phi(\mathbf{x}_i)) \leq 1](-y_i \phi(\mathbf{x}_i))$$

(ii) [3 points] You train a binary classifier by running stochastic gradient descent, and obtain a low training loss. However, you notice that the loss at test time is quite high. In 1-2 sentences, explain how you would change the hyperparameter λ to improve performance at test time and why this change would help.

Solution The model is overfitting to the training data. You should reduce the value of hyperparameter λ in order to increase regularization by increasing the relative weight of the $\frac{1}{2} \mathbf{w} \cdot \mathbf{w}$

(iii) [8 points] Consider a training dataset with two datapoints $(x_a, +1)$ and $(x_b, -1)$. You use a feature extractor $\phi(x)$ on these datapoints, and obtain $\phi(x_a) = [4 \ 2]^T$ and $\phi(x_b) = [1 \ 3]^T$.

You run stochastic gradient descent on $(x_a, +1)$, followed by $(x_b, -1)$ (in that order) with $\lambda = 0.1$ and $\eta = 0.1$. What is the value of the weight vector at the end of these two iterations? Show your work.

Solution Iteration 1 At iteration 1, $\mathbf{w} = \mathbf{0}$. Therefore, $y_a(\mathbf{w} \cdot \phi(\mathbf{x}_a)) = 0 < 1$. Therefore,

$$\nabla_{\mathbf{w}} L_i(\mathbf{w}) = \frac{1}{N} \mathbf{w} + \lambda \mathbf{1}[y_a(\mathbf{w} \cdot \phi(\mathbf{x}_a)) < 1](-y_a \phi(\mathbf{x}_a)) = -\lambda y_a \phi(\mathbf{x}_a)$$

Update rule is: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L_i(\mathbf{w}) = 0 - \eta(-\lambda y_a \phi(\mathbf{x}_a)) = \eta \lambda y_a \phi(\mathbf{x}_a) = 0.01 y_a \phi(\mathbf{x}_a)$

$$\mathbf{w} = 0.01 y_a \phi(\mathbf{x}_a) = [0.04 \ 0.02]^T$$

Iteration 2

At iteration 2, $y_b(\mathbf{w} \cdot \phi(\mathbf{x}_b)) = -0.1 < 1$. Therefore,

$$\nabla_{\mathbf{w}} L_i(\mathbf{w}) = \frac{1}{N} \mathbf{w} + \lambda \mathbf{1}[y_b(\mathbf{w} \cdot \phi(\mathbf{x}_b)) < 1](-y_b \phi(\mathbf{x}_b)) = \frac{1}{2} \mathbf{w} - \lambda y_b \phi(\mathbf{x}_b)$$

Update rule is: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L_i(\mathbf{w}) = \mathbf{w} - \eta(\frac{1}{2} \mathbf{w} - \lambda y_b \phi(\mathbf{x}_b)) = \mathbf{w} - 0.1(\frac{1}{2} \mathbf{w} - 0.1 y_b \phi(\mathbf{x}_b))$

$$\begin{aligned} \mathbf{w} &= [0.04 \ 0.02]^T - 0.1(\frac{1}{2}[0.04 \ 0.02]^T + 0.1[1 \ 3]^T) \\ &= [0.04 \ 0.02]^T - 0.1([0.12 \ 0.31]^T) \\ &= [0.04 \ 0.02]^T - [0.012 \ 0.031]^T \\ &= [0.028 \ -0.011]^T \end{aligned}$$

OR

$$\mathbf{w} = 0.001[28 \ -11]^T$$

b. (10 points) Modifying Margin

You now consider adding a hyperparameter $\hat{\delta} > 0$ to increase the minimum margin of the binary classifier. In other words, you define a modified pointwise loss function $\hat{Loss}(\mathbf{x}_i, y_i, \hat{\mathbf{w}})$ with hyperparameters $\hat{\lambda}$ and $\hat{\delta}$ as follows:

$$\hat{Loss}(\mathbf{x}_i, y_i, \hat{\mathbf{w}}) = \frac{1}{2N} \sum_{j=1}^d (\hat{w}^{(j)})^2 + \hat{\lambda} \max \left\{ 0, \hat{\delta} - y_i \left(\sum_{j=1}^d \hat{w}^{(j)} \phi(\mathbf{x}_i)^{(j)} \right) \right\} \quad \forall i \in \{1, 2, \dots, N\} \quad (2)$$

The weight vector $\hat{\mathbf{w}}$ is learned using stochastic gradient descent on the above loss function. Prove that $\hat{\delta}$ is an unnecessary hyperparameter, and our objective can be accomplished by tweaking the hyperparameter λ to some $\hat{\lambda}$. Specifically:

1. Prove the following by induction: for $\hat{\lambda} = \hat{\delta}\lambda$, show that $\hat{\mathbf{w}} = \hat{\delta}\mathbf{w}$ for iteration 1 and all subsequent iterations of the Stochastic Gradient Descent algorithm described in Algorithm 1. Note that \mathbf{w} and $\hat{\mathbf{w}}$ are both initialized to $\mathbf{0}$.
2. Prove that, for $\hat{\mathbf{w}} = \hat{\delta}\mathbf{w}$, weight vectors \mathbf{w} and $\hat{\mathbf{w}}$ will yield the same prediction labels on all data points \mathbf{x} . Note that predictions are given by $y_{pred}(\mathbf{x}; \mathbf{w}) = \text{sign}(\mathbf{w} \cdot \phi(\mathbf{x}))$

Solution Let $\hat{\lambda} = \hat{\delta}\lambda$.

We prove by induction that the weight vector $\hat{\mathbf{w}} = \hat{\delta}\mathbf{w}$ for iteration 1 and all subsequent iterations.

Let $\hat{\mathbf{w}}_t$ be the modified weight vector at iteration t .

Basis of proof: Prior to the first iteration (i.e., at iteration 0): $\hat{\mathbf{w}}_0 = \mathbf{0} = \hat{\delta}(\mathbf{0}) = \hat{\delta}\mathbf{w}_0$. Therefore, our premise holds true for iteration 0: $\hat{\mathbf{w}}_0 = \hat{\delta}\mathbf{w}_0$.

Alternative basis of proof - at iteration 1, using the update rule:

$$\hat{\mathbf{w}}_1 \leftarrow \hat{\mathbf{w}}_0 - \eta \nabla_{\hat{\mathbf{w}}_0} \hat{Loss}(\hat{\mathbf{w}}_0) = \mathbf{0} - \eta(-\hat{\lambda} y_i \phi(\mathbf{x}_i)) = \hat{\delta} \eta \lambda y_i \phi(\mathbf{x}_i) = \hat{\delta} \mathbf{w}_1$$

Therefore, our premise holds true for iteration 1: after the first update step, $\hat{\mathbf{w}}_1 = \hat{\delta}\mathbf{w}_1$. At subsequent update steps, we have the gradient calculated as follows:

$$\nabla_{\hat{\mathbf{w}}_t} \hat{Loss}(\hat{\mathbf{w}}_t) = \frac{1}{N} \hat{\mathbf{w}}_t + \hat{\lambda} \mathbf{1}[y_i(\hat{\mathbf{w}}_t \cdot \phi(\mathbf{x}_i)) < \hat{\delta}](-y_i \phi(\mathbf{x}_i))$$

Inductive step of proof: we assume that $\hat{\mathbf{w}}_t = \hat{\delta}\mathbf{w}_t$, and prove that $\hat{\mathbf{w}}_{t+1} = \hat{\delta}\mathbf{w}_{t+1}$

$$\begin{aligned} \hat{\mathbf{w}}_{t+1} &\leftarrow \hat{\mathbf{w}}_t - \eta \nabla_{\hat{\mathbf{w}}_t} \hat{Loss}(\hat{\mathbf{w}}_t) \\ &= \hat{\delta}\mathbf{w}_t - \eta \left(\frac{1}{N} \hat{\mathbf{w}}_t + \hat{\lambda} \mathbf{1}[y_i(\hat{\mathbf{w}}_t \cdot \phi(\mathbf{x}_i)) < \hat{\delta}](-y_i \phi(\mathbf{x}_i)) \right) \end{aligned}$$

$$\begin{aligned}
&= \hat{\delta}(\mathbf{w}_t - \eta(\frac{1}{N}\mathbf{w}_t + \hat{\lambda}\mathbf{1}[y_i(\mathbf{w}_t \cdot \phi(\mathbf{x}_i)) < 1](-y_i\phi(\mathbf{x}_i)))) \\
&= \hat{\delta}(\mathbf{w}_t - \eta\nabla_{\mathbf{w}_t} Loss(\mathbf{w}_t)) \\
&\Rightarrow \hat{\mathbf{w}}_{t+1} = \hat{\delta}(\mathbf{w}_{t+1})
\end{aligned}$$

This concludes the proof that $\hat{\mathbf{w}}_t = \hat{\delta}\mathbf{w}_t$ for all iterations $t \geq 1$.

Now, $y_{pred}(\mathbf{x}; \hat{\mathbf{w}}) = \text{sign}(\hat{\mathbf{w}} \cdot \phi(\mathbf{x})) = \text{sign}(\hat{\delta}\mathbf{w} \cdot \phi(\mathbf{x})) = \text{sign}(\mathbf{w} \cdot \phi(\mathbf{x})) = y_{pred}(\mathbf{x}; \mathbf{w})$ since $\hat{\delta} > 0$.

Therefore, vectors \mathbf{w} and $\hat{\mathbf{w}}$ yield the same predictions on all data points \mathbf{x} .

c. (6 points) Stopping Criteria

Our objective for this problem is to design stopping criteria that stop training once training loss has converged.

We define the average training loss at iteration \tilde{t} as follows:

$$L_{avg}(\mathbf{w})^{(\tilde{t})} = \frac{1}{\tilde{t}} \sum_{t=1}^{\tilde{t}} Loss(\mathbf{x}_i, y_i, \mathbf{w})^{(t)}$$

where $Loss(\mathbf{x}_i, y_i, \mathbf{w})^{(t)}$ is the pointwise loss obtained at iteration t .

We modify Algorithm 1 as follows (the only change to the algorithm is the calculation of average training loss at each iteration):

Algorithm 2: Modified Stochastic Gradient Descent

```
Randomly shuffle the training data
initialize  $\mathbf{w} = \mathbf{0}, i = 1, L_{avg} = 0$ 
while stopping criteria not reached do
  for  $t = 1, 2, \dots, T$  do
     $L_{avg} \leftarrow \frac{1}{t} (L_{avg} \times (t - 1) + Loss(\mathbf{x}_i, y_i, \mathbf{w}))$ 
     $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} Loss(\mathbf{x}_i, y_i, \mathbf{w})$ 
     $i \leftarrow (i \bmod N) + 1$ 
  end for
end while
```

(i) [3 points] Consider the following stopping criteria:

$$L_{avg}^{(\tilde{t})} < \epsilon \quad \text{and} \quad \tilde{t} \geq N$$

for some hyperparameter $\epsilon > 0$.

Is the criteria above good for our objective (i.e., designing stopping criteria that stop training once training loss has converged)? Explain why or why not.

Solution Our objective is to stop training when training loss converges (i.e., there is minimal change in the loss between iterations). The above criteria do not represent loss convergence, since they do not account for the change in loss between iterations. Loss could converge at a value above ϵ or may not have converged at a value below ϵ .

(ii) [3 points] Now consider the following stopping criteria:

$$\frac{|L_{avg}^{(\tilde{t})} - L_{avg}^{(\tilde{t}-1)}|}{L_{avg}^{(\tilde{t}-1)}} < \epsilon \quad \text{and} \quad \tilde{t} \geq N$$

for some hyperparameter $\epsilon > 0$. What issues might you run into by using the above criteria for stochastic gradient descent? (Hint: this new criteria might work well for regular gradient descent.)

Solution It does not account for the stochasticity in data points. The above criterion could stop training before training loss has actually converged. More specifically, all it takes to activate the above criteria is to find one point that yields a loss equal to average loss so far (within a margin of ϵ); this clearly does not mean that losses have actually converged.

Note that this could be a reasonable criterion for regular gradient descent, since one iteration of regular gradient descent goes through all points, so minimal change in loss between consecutive iterations actually indicates convergence.

2. Cluster Away (20 points)

Consider performing k -means clustering with $k = 2$ clusters on the following set of 3 one-dimensional data points:

$$\{-8, 0, 14\} \tag{3}$$

a. (4 points)

What is the globally optimal clustering and the associated reconstruction loss for the above dataset? Your answer should specify the following:

- Centroids μ_1 and μ_2 for clusters 1 and 2 respectively.
- Assignments of points $x_1 = -8, x_2 = 0$ and $x_3 = 14$ to clusters.
- Reconstruction loss (sum of squared distances) of this globally optimal clustering.

Solution $\mu_1, \mu_2 = -4, 14$

Assignments of points: $x_1 = -8$ and $x_2 = 0$ to cluster centered at -4 . $x_3 = 14$ to cluster centered at 14 .

Reconstruction loss: $(4)^2 + (4)^2 + 0^2 = \mathbf{32}$

b. (12 points)

Consider the following algorithm (henceforth referred to as **k-means++**) to initialize the k-means centroids intelligently rather than at random:¹

- I. Choose one centroid μ_1 , uniformly at random from among the data points.
- II. For each data point x , compute $D^*(x)$, the Euclidean distance between x and the nearest centroid that has already been chosen. In other words $D^*(x) = \min_{\mu_i} D(x, \mu_i)$
- III. Choose one new data point at random as a new centroid, using a weighted probability distribution where a point x is chosen with probability proportional to $D^*(x)$. In other words,

$$P[x_i \text{ is chosen as } \mu] = \frac{D^*(x_i)}{\sum_{j \in J} D^*(x_j)}$$

where J is the set of all points that haven't been chosen as a centroid yet.

- IV. Repeat Steps II and III until k centroids have been chosen.
- V. Now that the initial centroids have been chosen, proceed using standard k -means clustering.

Now consider the dataset specified in equation (3) with $k = 2$.

Fill in the table on the following page. For each row, you should specify (1) the probability of the centroids being chosen in this order with the k-means++ algorithm, (2) the **final** clusters that would result from these initial centroids (at the conclusion of k-means clustering), and (3) whether or not this clustering is optimal. You may leave your probabilities as unsimplified fractions and/or sums of fractions, e.g. $\frac{1}{1+2+3}$.

Note that, in the table on the following page, μ_1 is the first centroid chosen and μ_2 is the second centroid chosen.

¹Adapted from the following work: Arthur, David, and Sergei Vassilvitskii. k-means++: The advantages of careful seeding. Stanford, 2006.

μ_1	μ_2	Probability	Final Clusters	Optimal (Y/N)
-8	0			
-8	14			
0	-8			
0	14			
14	-8			
14	0			

What is the overall probability that using the k-means++ algorithm returns the globally optimal k -means cluster for this dataset with $k = 2$? (You may leave your answer as an unsimplified fraction or a sum of fractions.)

Solution

μ_1	μ_2	Probability	Final Clusters	Optimal (Y/N)
-8	0	$\frac{4}{45}$	$\{-8\}$ at centroid -8 and $\{0, 14\}$ at centroid 7	No
-8	14	$\frac{11}{45}$	$\{-8, 0\}$ at centroid -4 and $\{14\}$ at centroid 14	Yes
0	-8	$\frac{4}{33}$	$\{-8\}$ at centroid -8 and $\{0, 14\}$ at centroid 7	No
0	14	$\frac{7}{33}$	$\{-8, 0\}$ at centroid -4 and $\{14\}$ at centroid 14	Yes
14	-8	$\frac{11}{54}$	$\{-8, 0\}$ at centroid -4 and $\{14\}$ at centroid 14	Yes
14	0	$\frac{7}{54}$	$\{-8, 0\}$ at centroid -4 and $\{14\}$ at centroid 14	Yes

Probability of obtaining the optimal clustering = $\frac{11}{45} + \frac{7}{33} + \frac{11}{54} + \frac{7}{54} = \frac{391}{495} = 0.7899$

Detailed explanation of solution:

- Since the first centroid is assigned randomly,
 $Pr[\mu_1 = -8] = Pr[\mu_1 = 0] = Pr[\mu_1 = 14] = \frac{1}{3}$
- If $\mu_1 = -8$: $Pr[\mu_2 = 0|\mu_1 = -8] = \frac{8}{8+22} = \frac{4}{15}$
 $Pr[\mu_2 = 14|\mu_1 = -8] = \frac{22}{8+22} = \frac{11}{15}$
- if $\mu_1 = 0$: $Pr[\mu_2 = -8|\mu_1 = 0] = \frac{8}{8+14} = \frac{4}{11}$
 $Pr[\mu_2 = 14|\mu_1 = 0] = \frac{14}{8+14} = \frac{7}{11}$
- if $\mu_1 = 14$: $Pr[\mu_2 = -8|\mu_1 = 14] = \frac{22}{14+22} = \frac{11}{18}$
 $Pr[\mu_2 = 0|\mu_1 = 14] = \frac{14}{14+22} = \frac{7}{18}$

For initial centroids $\{-8, 0\}$, $x_3 = 14$ is assigned to cluster centered at $\mu = 0$ at iteration 1. Centroids are then updated to $\mu_1 = -8, \mu_2 = 7$. There is no change in assignments at the second iteration which leads to a final clustering of $\{-8\}$ at centroid $\mu_1 = -8$ and $\{0, 14\}$ at centroid $\mu_2 = 7$. This is suboptimal (optimal clustering is described in (a)).

For initial centroids either $\{-8, 14\}$ or $\{0, 14\}$, the optimal clustering described in (a) is obtained.

c. (4 points)

For part (c), consider an unspecified dataset in \mathbb{R}^d with n datapoints, and with $k \in \{2, 3, \dots, n\}$.

An alternative initialization method discussed in class is "multiple random initialization": (1) run multiple k-means clustering runs with random initializations of centroids on each run, and (2) choose the clustering with the lowest reconstruction loss.

Under what circumstances might you choose: (i) k-means++ initialization over multiple random initialization, (ii) multiple random initialization over k-means++ initialization? **Describe no more than one circumstance in each case.**

The circumstances you describe can include: the problem or application that clustering is being applied to, specific prior knowledge of the dataset, compute constraints, and/or the importance of getting the globally optimal clustering.

Solution We would choose k-means++ over multiple random initialization when compute is constrained, because multiple runs of k-means clustering is compute intensive.

We would choose multiple random initialization over k-means++ when compute costs are not all that important and accuracy is very highly valued (e.g., clustering to determine the likelihood of a malignant tumour or a similar medical application where getting an accurate clustering is very important). This is because a faulty k-means++ initialization - like any other one-time initialization method - could yield a poor (locally optimum but not globally optimum) clustering.

3. Weeding a lawn (*40 points*)

You are hired to weed a large lawn. Suppose the lawn is a $m \times n$ grid, and grass randomly grows in these grids (see Figure 1). Starting at a grid point, you turn on your weeding machine and weed the 3×3 area around you with no cost. Then, at each timestep, you must move to an adjacent cell and weed the 3×3 area centered at your new location at a total cost of 1 per timestep, until all grass is weeded. Near boundaries, you still weed the 3×3 area intersecting with the lawn (see Figure 2).

To formulate this as a state-based problem, each action is defined as a combination of move and weed, that is, a move from your current location followed by a weed of the 3×3 area centered at the location that you moved to, and you can assume there is no grass in the 3×3 area around the start location. Every move comes with cost 1.

You are expected to weed all the grass in the lawn. Let G_0 to be set of indices of all the grass in the initial lawn. For following questions, assume you start at $(1, 1)$ if not explicitly specified.




Figure 1: Example of a 4×5 lawn. Here, $G_0 = \{(4, 1), (4, 2), (1, 5), (3, 5)\}$. Standing on $(2, 3)$, if you move to $(2, 4)$, you can weed the entire red region.




Figure 2: Boundary example with the same lawn. Standing on $(1, 3)$, if you move to $(1, 4)$, you can weed the 6 squares in the intersection between the red region and the lawn.

a. (6 points)

Fill out the components of the search problem corresponding to the above problem setting. You can use the notation $G_{(i,j)}^{\text{weed}}$ to represent the set of locations of weeds that you can weed when moving to location (i, j) . You might find set difference notation useful: $A \setminus B = \{a : a \in A \text{ and } a \notin B\}$.

- $s_{\text{start}} = ((1, 1), G_0)$.
- $\text{Actions}(((i, j), G)) = \{a \in \{(-1, 0), (+1, 0), (0, -1), (0, +1)\} : (i, j) + a \text{ is in bounds}\}$.
- $\text{IsEnd}(((i, j), G)) =$

Solution $\text{IsEnd}(((i, j), G)) = \mathbf{1}[|G| = 0]$.

- $\text{Succ}(((i, j), G), a) =$

Solution $\text{Succ}(((i, j), G), a) = ((i, j) + a, G \setminus G_{(i,j)+a}^{\text{weed}})$

- $\text{Cost}(((i, j), G), a) =$

Solution $\text{Cost}(((i, j), G), a) = 1$.

b. (10 points)

You are very impatient for the original search problem to finish, so you want to apply A^* with "hierarchical pathfinding". To relax the problem, you divide the lawn into adjacent $k \times k$ blocks and in each move and weed action you weed all the grass in blocks that your 3×3 weeding area has intersection with. The original problem is equivalent to the relaxed problem with $k = 1$. **For simplicity, you can assume m and n are both multiples of k .**

Solving this relaxed search problem directly using UCS is as hard as the original problem. However, you can crack this relaxed problem using a modified UCS algorithm: when visiting a state, instead of moving one step, you will walk all the way to the border of next adjacent block with cost equal to the walked distance. Even though this algorithm only searches a subset of all possible paths, e.g. it will miss a zigzag path inside a block, it will always find a path with equal cost to any of optimal paths of the relaxed problem. In your answer you can assume the correctness of this algorithm and you don't need all details about this algorithm to solve following questions.

(1, 1)	(1, 2)	(1, 3)	(1, 4)	(1, 5)	(1, 6)	(1, 7)	(1, 8)
(2, 1)	(2, 2)	(2, 3)	(2, 4)	(2, 5)	(2, 6)	(2, 7)	(2, 8)
(3, 1)	(3, 2)	(3, 3)	(3, 4)	(3, 5)	(3, 6)	(3, 7)	(3, 8)
(4, 1)	(4, 2)	(4, 3)	(4, 4)	(4, 5)	(4, 6)	(4, 7)	(4, 8)
(5, 1)	(5, 2)	(5, 3)	(5, 4)	(5, 5)	(5, 6)	(5, 7)	(5, 8)
(6, 1)	(6, 2)	(6, 3)	(6, 4)	(6, 5)	(6, 6)	(6, 7)	(6, 8)
(7, 1)	(7, 2)	(7, 3)	(7, 4)	(7, 5)	(7, 6)	(7, 7)	(7, 8)
(8, 1)	(8, 2)	(8, 3)	(8, 4)	(8, 5)	(8, 6)	(8, 7)	(8, 8)

Figure 3: Example of a 8×8 lawn with block length $k = 4$. In the relaxed problem with the modified UCS algorithm, you can move to one of (1, 3), (2, 4), (4, 3) or (2, 1) with cost 1, 1, 2, 2 respectively. If you move to (4, 3), then your weeding area will intersect the lower-left gray block and you will weed the grass in (5, 1), (5, 4), and (7, 2).

(i) [6 points] Prove that we can use the future cost of this relaxed problem as a consistent heuristic for the original problem in part (a).

Solution The optimal path in original search problem is also a valid path in the relaxed search problem, meaning the optimal path in the relaxed problem must have equivalent or less cost. The modified UCS algorithm will find, from definition, one of the optimal path in the relaxed problem. These conclusions together guarantee that the future cost of this relaxed problem is less or equal to that of the original problem, thus a consistent heuristic.

(ii) [4 points] What is one advantage for using large k for this relaxed problem? What is one advantage of using small k ? Please limit your answer to a maximum of two sentences for each question.

Solution Large k : easier relaxed problem. Small k : more accurate heuristics.

c. (12 points)

Suppose you are standing at (3, 4) and you've managed to weed the entire lawn! Before you celebrate your success, however, you see that grass is about to re-grow in exactly one of three possible places: (1, 1), (4, 1), and (2, 9) *with uniform probability*. You can see this in Figure 4. You need to make one action before they re-grow, but you don't know which spot the grass will pop up in! After this action you will know the location of the newly grown grass and head to weeding it.





(1, 1) 	(1, 2)	(1, 3)	(1, 4)	(1, 5)	(1, 6)	(1, 7)	(1, 8)	(1, 9)
(2, 1)	(2, 2)	(2, 3)	(2, 4)	(2, 5)	(2, 6)	(2, 7)	(2, 8)	(2, 9) 
(3, 1)	(3, 2)	(3, 3)	(3, 4) 	(3, 5)	(3, 6)	(3, 7)	(3, 8)	(3, 9)
(4, 1) 	(4, 2)	(4, 3)	(4, 4)	(4, 5)	(4, 6)	(4, 7)	(4, 8)	(4, 9)

Figure 4: Lawn of this problem. You stand on (3, 4) and grass will grow in exactly one of the three marked locations with uniform probability after your next move.

(i) [6 points] What is the optimal action you can take to minimize the *expected future cost*? If there is a tie between optimal actions, state all of them. What is the expected future cost *after* you take this action?

Solution Move left (0, -1), because moving down/right will bring you closer to one grass while farther from other two, increasing the total cost in expectation; moving up will bring you closer to (1, 1) but farther from (4, 1), keeping total expected cost same. The expected cost for moving left would be $\frac{1}{3}(1 + 2 + 5) = \frac{8}{3}$.

(ii) [6 points] What is the optimal action you can take to minimize the *worst case future cost*? If there is a tie between optimal actions, state all of them. What is the worst case future cost *after* you take this action?

Solution Move up (-1, 0)/right (0, 1), because your distances to three grass are 3, 2 and 4. Moving either up/right will bring the maximum distance to 4. For example, if you move right, the distances to (1, 1), (4, 1) and (2, 9) are 4, 3 and 3 respectively with the maximum 4. The worst case cost would be 4.

d. (12 points)

You find that grass randomly grows in certain locations on the lawn, denoted in the set H . More specifically, after every move and weed you make, grass will regrow in one location from H , chosen uniformly at random. If there is already grass at that location, nothing will happen. Grass can also re-grow at the location that you have just weeded.

Because it is now impossible to weed all the grass, you are paid at a per-piece basis: you will receive reward q for each grass you weed, and you can quit at any time. The moving cost 1 (or equivalently reward of -1) still applies to this problem.

Due to the randomness of regrown grass, taking an action might transit to a state indeterminately, so you want to form it as a MDP problem instead.

(i) [6 points] Modify the original search problem into a MDP. You can use the notation $G_{(i,j)}^{\text{weed}}$ to represent the set of locations of weeds that you can weed when moving to location (i, j) .

- $s_{\text{start}} = ((1, 1), G_0)$.
- $\text{Actions}(((i, j), G)) = \{a \in \{(-1, 0), (+1, 0), (0, -1), (0, +1)\} : (i, j) + a \text{ is in bounds}\} \cup \{\text{Quit}\}$
- For $a \neq \text{Quit}$, $\text{Reward}(((i, j), G), a, ((i, j) + a, G')) =$

Solution $\text{Reward}(((i, j), G), a, ((i, j) + a, G')) = |G_{(i,j)+a}^{\text{weed}} \cap G| \times q - 1$

Some students interpreted the notation $G_{(i,j)}^{\text{weed}}$ as only containing the locations with weeds, instead of all locations in the square that could be weeded. In this case, we accepted an alternate solution for the reward:

$$\text{Reward}(((i, j), G), a, ((i, j) + a, G')) = |G_{(i,j)+a}^{\text{weed}}| \times q - 1$$

- For $a \neq \text{Quit}$, $T(((i, j) + a, G') \mid ((i, j), G), a) =$

Solution $T(((i, j) + a, G') \mid ((i, j), G), a) =$

$$\begin{cases} \frac{1}{|H|} & \text{if } \exists h \in H, G' = G \setminus G_{(i,j)+a}^{\text{weed}} \cup \{h\} \text{ and } h \notin G \setminus G_{(i,j)+a}^{\text{weed}} \text{ (or } G' \neq G \setminus G_{(i,j)+a}^{\text{weed}}) \\ \frac{|(G \setminus G_{(i,j)+a}^{\text{weed}}) \cap H|}{|H|} & \text{if } G' = G \setminus G_{(i,j)+a}^{\text{weed}} \\ 0 & \text{otherwise} \end{cases}$$

The first case denotes the case where we have a grass in G' that we got from regrowing in a spot that was not in G . The second case is if a grass tried to regrow in a spot that

was already in G , so nothing happened and therefore G' matches G . The third case is for any cases that will not match up with our successor state G' .

Note that for case 2, G' looks the same as G ; therefore, having a new grass try to regrow in any location $h \in H$ that is also in G will end up with the same successor state (as opposed to case 1, where G' might look different for different locations $h \in H$).

(ii) [6 points] You are thinking about using value iteration on this MDP, and you are interested in its time complexity.

- For any state, what is the largest possible number of successor states? Provide an exact expression.
- What is the total number of states possible for this problem, using Big-O notation?
- What is the time complexity of one iteration of value iteration of all states, using Big-O notation?

For all of these questions, you may express your answer using m , n , G , and H , though you may not need all of these.

Solution

- For any state, the largest possible number of successor states is $4|H|+1$ states, including {Quit}. We can break down each action into move and grow. Move takes 4 possible actions and gives deterministic result while grow is a single action that transitions to at most $|H|$ other states.

Solutions that did not consider a quitting state and just gave $4|H|$ were also accepted.

- Since the state space is the Cartesian product of possible x axis values, possible y axis values and possible subsets of all grass locations, the number of states = $O(mn2^{|G \cup H|})$. $O(mn2^{mn})$ is also a valid bound, though not the tightest.
- Therefore the time complexity of one iteration of value iteration for all states is the product of the two previous values, or $O(|H|mn2^{|G \cup H|})$.

4. Olympics (50 points)

You are the president of the small nation of Inferencia, and you have been charged with choosing which of your country's two rival soccer teams - the Bayesians or the Markovians - should represent Inferencia at the upcoming Olympics. You'd like to send whichever team is more popular, so you decide to model the monthly evolution of the two teams' fanbases during the months leading up to the Olympics using a dynamic Bayesian network.

Let B_t denote the number of fans that the Bayesians have in month t , and let M_t denote the number of fans that the Markovians have in month t . You have no way of observing these quantities directly, but you can observe two other quantities which they influence: let J_t denote the number of jerseys sold by the Bayesians in month t , and let A_t denote the attendance of the monthly exhibition game between the Bayesians and the Markovians in month t .

The fanbases of the two teams evolve according to the following model, where each month a fan is either gained or lost with equal probability:

$$\Pr(M_{t+1}|M_t) = \begin{cases} \frac{1}{2} & \text{if } M_{t+1} = M_t - 1 \\ \frac{1}{2} & \text{if } M_{t+1} = M_t + 1 \\ 0 & \text{otherwise} \end{cases} \quad \Pr(B_{t+1}|B_t) = \begin{cases} \frac{1}{2} & \text{if } B_{t+1} = B_t - 1 \\ \frac{1}{2} & \text{if } B_{t+1} = B_t + 1 \\ 0 & \text{otherwise} \end{cases}$$

The Bayesian fans are big spenders - almost every fan buys a jersey each month! We model the fanbase size's influence on jersey sales by:

$$\Pr(J_t|B_t) = \begin{cases} 0.3 & \text{if } J_t = B_t \\ 0.25 & \text{if } J_t = B_t - 1 \\ 0.2 & \text{if } J_t = B_t - 2 \\ 0.15 & \text{if } J_t = B_t - 3 \\ 0.1 & \text{if } J_t = B_t - 4 \\ 0 & \text{otherwise} \end{cases}$$

Lastly, because most fans attend each monthly exhibition (although sometimes more, and sometimes fewer), we model the influence of the fanbase sizes on the exhibition attendance by:

$$\Pr(A_t|B_t, M_t) = \begin{cases} 0.14 & \text{if } A_t = B_t + M_t \\ 0.13 & \text{if } |A_t - (B_t + M_t)| = 1 \\ 0.11 & \text{if } |A_t - (B_t + M_t)| = 2 \\ 0.09 & \text{if } |A_t - (B_t + M_t)| = 3 \\ 0.06 & \text{if } |A_t - (B_t + M_t)| = 4 \\ 0.04 & \text{if } |A_t - (B_t + M_t)| = 5 \\ 0 & \text{otherwise} \end{cases}$$




Figure 5: The changing fanbases process modeled as a dynamic Bayesian network. The unshaded nodes correspond to the latent/hidden fanbase counts, and the shaded nodes correspond to the observable emissions.

Note that the assumptions and inferences made in individual parts (i.e. **(a)**, **(b)**, etc.) of this problem do *not* carry over from one to the next; the only assumptions you may make in a given part are those which are explicitly stated in that part's description.

a. (6 points) (Conditional) Independences

Mark each of the following as True or False.

(i) [1 point] $B_t \perp\!\!\!\perp J_{t+1}$

Solution False

(ii) [1 point] $B_t \perp\!\!\!\perp J_{t+1} \mid B_{t+1}$

Solution True

(iii) [1 point] $B_t \perp\!\!\!\perp M_t$

Solution True

(iv) [1 point] $B_t \perp\!\!\!\perp M_t \mid A_t$

Solution False

(v) [1 point] $A_t \perp\!\!\!\perp M_{t+1}$

Solution False

(vi) [1 point] $A_t \perp\!\!\!\perp M_{t+1} \mid M_t$

Solution True

b. (8 points) Domain Consistencies

As a first step, we will not concern ourselves with which fanbase counts are *probable*, but instead which counts are even *possible*. Suppose that we observe, in our first month of collecting data, that $J_1 = 75$ and $A_1 = 100$. Give the domains for M_1 and B_1 that are consistent with these observations.

You need only give the consistent domains (using either set notation or inequality notation) in order to receive full credit.

Solution The only values of B_1 which are consistent (i.e. yield nonzero probability under $\Pr(J_1|B_1)$) with $J_1 = 75$ are $B_1 \in \{75, 76, 77, 78, 79\}$.

Now, we can use this reduced domain for B_1 together with the fact that $A_1 = 100$ to reason about the domain of M_1 . To yield nonzero probability under $\Pr(A_1|B_1, M_1)$, we must have:

$$|A_1 - (B_1 + M_1)| = |100 - (B_1 + M_1)| \leq 5$$

And to satisfy this inequality, we must have:

$$95 \leq B_1 + M_1 \leq 105$$

If $B_1 = 75$, we see that $20 \leq M_1 \leq 30$. Because M_1 is largest when B_1 is smallest, this gives an upper bound of 30 on the domain of M_1 . Similarly, if $B_1 = 79$, we see that $16 \leq M_1 \leq 26$. Because M_1 is smallest when B_1 is largest, this gives a lower bound of 16 on the domain of M_1 . Thus, we conclude that $16 \leq M_1 \leq 30$.

c. (10 points) Inference

Suppose the Bayesian's manager took a nationwide poll in month t that concluded they had exactly 75 fans. Suppose additionally that in month $t + 2$, the Bayesians sell 73 jerseys. What is the probability that in month $t + 2$ the Bayesians have 77 fans?

$$\Pr(B_{t+2} = 77|B_t = 75, J_{t+2} = 73) =$$

Solution By Bayes rule, we have:

$$\Pr(B_{t+2} = 77|B_t = 75, J_{t+2} = 73) = \frac{\Pr(J_{t+2} = 73|B_t = 75, B_{t+2} = 77)\Pr(B_{t+2} = 77|B_t = 75)}{\Pr(J_{t+2} = 73|B_t = 75)}$$

We'll begin with the first term in the numerator; because J_{t+2} is conditionally independent of B_t given B_{t+2} , we have $\Pr(J_{t+2} = 73|B_t = 75, B_{t+2} = 77) = \Pr(J_{t+2} = 73|B_{t+2} = 77)$. This is simply given by our jersey sales model; the probability that the Bayesians sell four fewer jerseys than they have fans is 0.1.

We turn next to the second term in the numerator; if there are 75 fans in month t , then with equal probability there are either 74 or 76 fans in month $t + 1$. If there were 74 in month $t + 1$, then there would be either 73 or 75 in month $t + 2$ with equal probability, and if there were 76 in month $t + 1$, then there would be either 75 or 77 in month $t + 2$ with equal probability. Thus, we have that $\Pr(B_{t+2} = 73|B_t = 75) = \Pr(B_{t+2} = 77|B_t = 75) = 0.25$, and $\Pr(B_{t+2} = 75|B_t = 75) = 0.5$.

Now, to compute the denominator, we simply sum the expression in the numerator across all possible values for B_{t+2} :

$$\Pr(J_{t+2} = 73|B_t = 75) = \sum_x \Pr(J_{t+2} = 73|B_t = 75, B_{t+2} = x)\Pr(B_{t+2} = x|B_t = 75)$$

Following the same reasoning as we used for the numerator, this evaluates to:

$$\begin{aligned} &= 0.25 \cdot \Pr(J_{t+2} = 73|B_{t+2} = 73) + 0.5 \cdot \Pr(J_{t+2} = 73|B_{t+2} = 75) + 0.25 \cdot \Pr(J_{t+2} = 73|B_{t+2} = 77) \\ &= 0.25 \cdot 0.3 + 0.5 \cdot 0.2 + 0.25 \cdot 0.1 = 0.075 + 0.1 + 0.025 = 0.2 \end{aligned}$$

So altogether, we have:

$$\Pr(B_{t+2} = 77|B_t = 75, J_{t+2} = 73) = \frac{0.1 \cdot 0.25}{0.2} = \frac{1}{2} \cdot \frac{1}{4} = \frac{1}{8}$$

d. (4 points) Gibbs Sampling

Inference is exhausting; you decide that you'd be satisfied with simply being able to draw samples from distributions rather than specifying them exactly. In particular, you want to sample joint assignments to the variables $\{B_t, M_t, A_t, J_t\}_{t=1}^T$ for some time horizon T . You decide to implement Gibbs sampling for this purpose, but something's not right! What additional information, beyond what we've given you, would allow you to perform Gibbs sampling? Briefly explain.

Solution (The following argument applies identically to M_t as well as B_t): In order to sample B_t , we need to have first assigned a value to B_{t-1} ; but in order to have sampled a value for B_{t-1} , we need to have first assigned a value to B_{t-2} , and so on. Continuing in this way, we realize that we must have a way of assigning a value to B_1 in order to perform Gibbs sampling. But to do this, we would either need to specify a fixed value for B_1 , or specify a prior distribution $\Pr(B_1)$ from which to sample.

e. (12 points) Exact Filtering

You now want to begin making inferences as to the sizes of the teams' fanbases given only observations of attendances and jersey sales. Recall that exact inference of this kind in dynamic Bayesian networks can be achieved using a dynamic programming approach - for example, in the context of Hidden Markov Models, we used the forward-backward algorithm to do filtering and smoothing.

Give recursive expressions for the following filtering queries. Leave your expressions in terms of known probabilities.

(i) [4 points] Let's start by making inferences based only on observed jersey sales. Denote $F_t(b_t) = \Pr(B_t = b_t | J_1 = j_1, \dots, J_t = j_t)$. Give a recursive expression for $F_t(b_t)$ assuming that you've already computed $F_{t-1}(b_{t-1})$ for all b_{t-1} .

Solution This is exactly the “forward” computation in an HMM. We can compute the unnormalized quantity, which we'll denote $\tilde{F}_t(b_t)$, using the standard forward update:

$$\tilde{F}_t(b_t) = \sum_{b_{t-1}} F_{t-1}(b_{t-1}) \cdot \Pr(B_t = b_t | B_{t-1} = b_{t-1}) \cdot \Pr(J_t = j_t | B_t = b_t)$$

and can subsequently produce the required probability by normalizing:

$$F_t(b_t) = \frac{\tilde{F}_t(b_t)}{\sum_{b'_t} \tilde{F}_t(b'_t)}$$

(ii) [8 points] Let's bring in the observed attendances as well! Now, denote

$F_t(b_t, m_t) = \Pr(B_t = b_t, M_t = m_t | J_1 = j_1, \dots, J_t = j_t, A_1 = a_1, \dots, A_t = a_t)$. Give a recursive expression for $F_t(b_t, m_t)$ assuming that you've already computed $F_{t-1}(b_{t-1}, m_{t-1})$ for all b_{t-1} and all m_{t-1} .

Solution This closely mirrors the “forward” computation in an HMM, but now we must account for the dynamics of both hidden states, as well as the probabilities of both observed emissions. We can compute the unnormalized quantity, which we'll denote $\tilde{F}_t(b_t, m_t)$, using the following forward update:

$$\begin{aligned} \tilde{F}_t(b_t, m_t) = \\ \sum_{b_{t-1}, m_{t-1}} F_{t-1}(b_{t-1}, m_{t-1}) \cdot \Pr(B_t = b_t | B_{t-1} = b_{t-1}) \cdot \Pr(M_t = m_t | M_{t-1} = m_{t-1}) \cdot \\ \Pr(J_t = j_t | B_t = b_t) \cdot \Pr(A_t = a_t | B_t = b_t, M_t = m_t) \end{aligned}$$

and can subsequently produce the required probability by normalizing:

$$F_t(b_t, m_t) = \frac{\tilde{F}_t(b_t, m_t)}{\sum_{b'_t, m'_t} \tilde{F}_t(b'_t, m'_t)}$$

f. (10 points) Particle Filtering

Throughout this problem, you are free to leave quantities in terms of unevaluated expressions (i.e. you may write $0.75 \cdot 0.5$ instead of 0.375).

Computing all of those terms exactly seems tedious, so you instead decide to employ particle filtering to quickly and painlessly provide you with approximate solutions. You're fine with a (very) crude approximation, so you only use two particles.

(i) [2 points] Suppose you begin with the two particles $(B_1 = 80, M_1 = 75)$ and $(B_1 = 82, M_1 = 74)$. You then observe that $J_1 = 79$ and $A_1 = 154$. Compute the weights that you should assign to the two particles based on this evidence.

Solution For the first particle, we have $\Pr(A_1 = 154|B_1 = 80, M_1 = 75) = 0.13$ and $\Pr(J_1 = 79|B_1 = 80) = 0.25$. Thus, the first particle should get a weight of $0.13 * 0.25 = 0.0325$.

Similarly, for the second particle, we have $\Pr(A_1 = 154|B_1 = 82, M_1 = 74) = 0.11$ and $\Pr(J_1 = 79|B_1 = 82) = 0.15$. Thus, the second particle should get a weight of $0.11 * 0.15 = 0.0165$.

(ii) [2 points] Using these weights, we now resample two new particles. Provide this sampling distribution.

Probability of sampling a new particle to be $(B_1 = 80, M_1 = 75) =$

Solution $\frac{0.0325}{0.0325+0.0165}$

Probability of sampling a new particle to be $(B_1 = 82, M_1 = 74) =$

Solution $\frac{0.0165}{0.0325+0.0165}$

(iii) [3 points] Suppose both of our new particles are sampled to be $(B_1 = 80, M_1 = 75)$. We now extend these particles using our dynamics models. What is the probability that a particular one of these two particles is extended to:

$$(B_1 = 80, M_1 = 75, B_2 = 78, M_2 = 76)?$$

Solution Zero. Under the given model for $\Pr(B_{t+1}|B_t)$, the only possible values for B_2 are 79 and 81.

$$(B_1 = 80, M_1 = 76, B_2 = 79, M_2 = 75)?$$

Solution Zero. The value assigned to M_1 cannot change upon extending the particle.

$$(B_1 = 80, M_1 = 75, B_2 = 79, M_2 = 76)?$$

Solution $\Pr(B_2 = 79|B_1 = 80) \cdot \Pr(M_2 = 76|M_1 = 75) = \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{4}$

(iv) [3 points] Suppose now that you have access to a large number of particles which are approximating the distribution over $(B_1, \dots, B_n, M_1, \dots, M_n)$. The Olympics are happening in 6 months, but you have to decide now which team to send so that they can start preparing! You decide to make predictions of B_{n+6} and M_{n+6} in order to send whichever team you predict to be more popular during the month in which the Olympics will be held. Explain in a few sentences how you would use your particles for making this decision.

Solution Propagate each particle through the two dynamics models six times in order to sample values of B_{n+1}, \dots, B_{n+6} and M_{n+1}, \dots, M_{n+6} for each particle. Compute the average values of B_{n+6} and M_{n+6} across all of the particles, and send whichever team has the larger average value.